

Experiment: 1 Simple Artificial neural network

Aim:

To build a simple Artificial Neural Network (ANN) classification using customer churn dataset.

Software Requirements:

➤ Google Colab

Program:

#Import Libraries

```
import numpy as np
import pandas as pd
import tensorflow as tf
```

```
from google.colab import drive
drive.mount('/content/drive')
```

Load dataset

```
data = pd.read_csv('Churn_Modelling.csv')
data
```

Extract features and label

```
X = data.iloc[:, 3:-1].values
print(X)
Y = data.iloc[:, -1].values
print(Y)
```

Encode categorical data

```
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
X[:, 2] = np.array(le.fit_transform(X[:, 2]))
```

```
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder
ct = ColumnTransformer(transformers=[('encoder', OneHotEncoder(), [1])],
remainder='passthrough')
X = np.array(ct.fit_transform(X))
```

Split dataset into training and testing sets

```
from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=0)
```

Feature scaling

```
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

Build the ANN model

```
ann = tf.keras.models.Sequential()
ann.add(tf.keras.layers.Dense(units=6, activation='relu'))
ann.add(tf.keras.layers.Dense(units=1, activation='sigmoid'))
```

Compile the model

```
ann.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

Train the model

```
ann.fit(X_train, Y_train, batch_size=32, epochs=100)
```

Predict a new result

```
print(ann.predict(sc.transform([[0,1,1,619,0,42,2,60000,1,1,1,101348]])) > 0.5)
```

→ Output

```
[[ True]]
```

Result:

The experiment successfully resulted in a Simple artificial neural network capable of predicting customer churn, achieving accurate classification .

Experiment: 2 Single Layer Perceptron

Aim:

To implement a Single Layer Perceptron using the Iris dataset for binary classification and evaluate its performance.

Software Requirements:

- Google Colab

Program:

#Import Libraries

```
from sklearn.datasets import load_iris
from sklearn.linear_model import Perceptron
from sklearn.metrics import accuracy_score
```

Load the Iris dataset

```
iris = load_iris()
```

Select features and target

```
x = iris.data[:, (2, 3)] # Petal length and petal width
y = (iris.target == 0).astype(int) # Binary classification: Setosa or not
```

Initialize the Perceptron

```
ptron = Perceptron(random_state=42)
```

Train the model

```
ptron.fit(x, y)
```

Make predictions

```
y_pred = ptron.predict(x)
print(y_pred)
```

Evaluate the model

```
print(f'Accuracy Score: {accuracy_score(y, y_pred)}')
```

→ Output

```
Accuracy Score: 1.0
```

Result :

The experiment successfully demonstrated the working of a Single Layer Perceptron on the Iris dataset, achieving accurate binary classification

Experiment: 3 Gradient Descent

Aim:

To implement the Gradient Descent optimization algorithm for minimizing Errors

Software Requirements:

➤ Google Colab

Program:

#Import Libraries

```
import numpy as np
```

Gradient Descent Function

```
def gradient_descent(gradient, start, learning_rate, iteration=50, tol=1e-06):
```

```
    vector = start
```

```
    for _ in range(iteration):
```

```
        diff = -learning_rate * gradient(vector)
```

```
        if np.all(np.abs(diff) <= tol):
```

```
            break
```

```
        vector += diff
```

```
    return vector
```

Testing The Function

```
from typing_extensions import LiteralString
```

```
print(gradient_descent(gradient = lambda v: 4* v**3 - 10* v-3,start=0,learning_rate=0.2 ))
```

Result:

The experiment successfully demonstrated the use of the Gradient Descent algorithm

➔ Output

```
-1.4207567437458342
```

Experiment: 4 Stochastic Gradient Descent-Classifier

Aim:

To implement and evaluate a linear classification model using Stochastic Gradient Descent (SGD) with elastic net regularization.

Software Requirements:

- Google Colab

Program:

```
import numpy as np
```

```
from sklearn import linear_model
```

Sample dataset

```
x = np.array([[-1, -1], [-2, -1], [1, 1], [2, 1]])
```

```
y = np.array([1, 1, 2, 2])
```

Creating SGDClassifier model with elasticnet regularization

```
sdg_class = linear_model.SGDClassifier(max_iter=1000, tol=1e-3, penalty="elasticnet")
```

Fitting model on the original data

```
sdg_class.fit(x, y)
```

Making prediction

```
print('Prediction is:', sgd_class.predict([[2, 4]]))
```

Model parameters

```
print('Weight vector(s):', sgd_class.coef_)
```

```
print('Intercept:', sgd_class.intercept_)
```

```
print('Distance to HyperPlane:', sgd_class.decision_function([[2, 4]]))
```

→ Output

```
Prediction is : [2]
```

```
Weight vector(s): [[9.77200712 9.77200712]]
```

```
Intercept: [-10.]
```

```
Distance to HyperPlane : [48.63204273]
```

Result:

The Stochastic Gradient Descent classifier is able to make predictions and display the model's weight and intercept values and successfully verified.

Experiment: 5 Fundamentals Of TensorFlow

Aim:

To understand and perform basic operations in TensorFlow, including constants, variables, reshaping, matrix multiplication, and concatenation.

Software Requirements:

- Google Colab

Program:

```
import tensorflow as tf  
print("TensorFlow version:", tf.__version__)
```

→ **Output:** `TensorFlow version: 2.18.0`

Check if Eager Execution is enabled

```
if(tf.executing_eagerly()):  
    print("Eager Execution Enabled")  
else:  
    print("Eager Execution Not Available. Upgrade TensorFlow 2.0.0+")
```

→ **Output:** `Eager Execution Enabled`

Constants

```
con1 = tf.constant([[1.4, 2.1], [3, 4.7]])  
con2 = tf.constant([[5], [2]])  
con3 = tf.constant([[5, 4], [2, 1]])  
con4 = tf.constant([[5, 4, 2], [2, 1, 2]])  
print("T1:", con1)  
print("T2:", con2)  
print("T3:", con3)
```

→ **Output:** T1: tf.Tensor(
[[1.4 2.1]
[3. 4.7]], shape=(2, 2), dtype=float32)
T2: tf.Tensor(
[[5]
[2]], shape=(2, 1), dtype=int32)
T3: tf.Tensor(
[[5 4]
[2 1]], shape=(2, 2), dtype=int32)

String tensor

```
T1 = tf.constant([[ "a"], [ "b"]], dtype=tf.string)
print(T1)
```

→ **Output:** <tf.Tensor: shape=(2, 1), dtype=string, numpy=
array([[b'a'],
[b'b']], dtype=object)>

Transpose

```
trans = tf.transpose(con1)
print("con1 Transpose:", trans)
```

→ **Output:** con1 Transpose: tf.Tensor(
[[1.4 3.]
[2.1 4.7]], shape=(2, 2), dtype=float32)

Type casting

```
con3 = tf.cast(con1, tf.float32)
con4 = tf.cast(con2, tf.float32)
```

→ **Output:** <tf.Tensor: shape=(2, 1), dtype=float32, numpy=
array([[5.],
[2.]], dtype=float32)>

Element-wise multiplication

```
mul_elements = tf.multiply(con3, con4)
print("Mul Elements:", mul_elements)
```

→ **Output:** Mul Elements: tf.Tensor(
[[7. 10.5]
[6. 9.4]], shape=(2, 2), dtype=float32)

Matrix multiplication

```
mul_mat = tf.matmul(con3, con4)
```

```
print("Mul Matrix:", mul_mat)
```

→ **Output:** Mul Matrix: tf.Tensor(
[[11.2]
[24.4]], shape=(2, 1), dtype=float32)

Reshaping

```
reshape_con1 = tf.reshape(tensor=con1, shape=[1, 4])
```

```
print("Reshape Con1:", reshape_con1)
```

→ **Output:** Reshape Con1: tf.Tensor([[1.4 2.1 3. 4.7]], shape=(1, 4),
dtype=float32)

Create a 6x6 tensor matrix

```
con7 = tf.constant([ [1.4, 2.1, 2.7, 2.8, 2.7, 8.7], [3.1, 4.7, 5.5, 1.4, 2.1, 2.7],  
                    [4.5, 2.5, 3.1, 4.7, 5.5, 5.6], [4.2, 2.2, 4.5, 2.5, 3.1, 4.7],  
                    [6.5, 7.5, 2.5, 3.1, 4.7, 5.5], [8.2, 9.2, 4.5, 2.5, 3.1, 4.7]])
```

```
print(con7)
```

→ **Output:** tf.Tensor([[1.4 2.1 2.7 2.8 2.7 8.7]
[3.1 4.7 5.5 1.4 2.1 2.7]
[4.5 2.5 3.1 4.7 5.5 5.6]
[4.2 2.2 4.5 2.5 3.1 4.7]
[6.5 7.5 2.5 3.1 4.7 5.5]
[8.2 9.2 4.5 2.5 3.1 4.7]], shape=(6, 6), dtype=float32)

Reshape the matrix with Total Element

```
reshape_con7 = tf.reshape(tensor=con7, shape=[6, 6])
```

```
print("Reshape con7:", reshape_con7)
```

→ **Output:** Reshape Con1: tf.Tensor(
[[1.4 2.1 2.7 2.8 2.7 8.7]
[3.1 4.7 5.5 1.4 2.1 2.7]
[4.5 2.5 3.1 4.7 5.5 5.6]
[4.2 2.2 4.5 2.5 3.1 4.7]
[6.5 7.5 2.5 3.1 4.7 5.5]
[8.2 9.2 4.5 2.5 3.1 4.7]], shape=(6, 6), dtype=float32)

Identity matrix

```
id_int = tf.eye(num_rows=3, num_columns=3, dtype=tf.int32)
```

```
print("Identity matrix (int):", id_int)
```

→ **Output:** identity matrix: tf.Tensor(
[[1 0 0]
[0 1 0]
[0 0 1]], shape=(3, 3), dtype=int32)

```
id_float = tf.eye(num_rows=3, num_columns=3, dtype=tf.float32)
```

```
print("Identity matrix (float):", id_float)
```

→ **Output:** identity matrix: tf.Tensor(
[[1. 0. 0.]
[0. 1. 0.]
[0. 0. 1.]], shape=(3, 3), dtype=float32)

Constant tensor

```
con_ten = tf.constant([[4, 1], [3, 4]])
```

```
print(con_ten)
```

→ **Output:** tf.Tensor(
[[4 1]
[3 4]], shape=(2, 2), dtype=int32)

Variable tensor

```
new_var_ten = tf.Variable([[4, 1], [3, 4]])
```

```
print(new_var_ten)
```

→ **Output:** tf.Tensor(
[[4 1]
[3 4]], shape=(2, 2), dtype=int32)

```
n_var_ten = new_var_ten.assign([[4, 1], [3, 4]])
```

```
print(n_var_ten)
```

→ **Output:** <tf.Variable 'UnreadVariable' shape=(2, 2) dtype=int32, numpy=
array([[4, 1],
[3, 4]], dtype=int32)>

Concatenation

```
row_con = tf.concat(values=(con_ten, new_var_ten), axis=0)

print(row_con)
```

→ **Output:** `tf.Tensor(`
`[[4 1]`
`[3 4]`
`[4 1]`
`[3 4]], shape=(4, 2), dtype=int32)`

```
col_con = tf.concat(values=(n_var_ten, new_var_ten), axis=1)

print(col_con)
```

→ **Output:** `tf.Tensor(`
`[[4 1 4 1]`
`[3 4 3 4]], shape=(2, 4), dtype=int32)`

Zeros and Ones

```
zeros = tf.zeros(shape=(3, 4), dtype=tf.int32)

print(zeros)
```

→ **Output:** `tf.Tensor(`
`[[0 0 0 0]`
`[0 0 0 0]`
`[0 0 0 0]], shape=(3, 4), dtype=int32)`

```
ones = tf.ones(shape=(3, 4), dtype=tf.int32)

print(ones)
```

→ **Output:** `tf.Tensor(`
`[[1 1 1 1]`
`[1 1 1 1]`
`[1 1 1 1]], shape=(3, 4), dtype=int32)`

Result:

The experiment successfully demonstrated various fundamental operations in TensorFlow including tensor creation, reshaping, matrix operations, concatenation, and use of constants and variables.

Experiment: 6 Working with Keras

Aim:

To understand and explore the basic functionalities of the Keras library, including loading datasets, visualizing images, text vectorization, and normalization.

Software Requirements:

- Google Colab

Program:

Importing libraries

```
from tensorflow import keras
```

```
import numpy as np
```

```
import tensorflow as tf
```

```
from matplotlib import pyplot as plt
```

```
from keras.datasets import cifar10, mnist, fashion_mnist
```

```
from tensorflow.keras.layers import TextVectorization, Normalization
```

Load and visualize CIFAR-10

```
tf.keras.datasets
```

```
data = cifar10.load_data()
```

#Data Split

```
(trainX, trainY), (testX, testY) = cifar10.load_data()
```

```
print('Train: X=%s, y=%s' % (trainX.shape, trainY.shape))
```

```
print('Test: X=%s, y=%s' % (testX.shape, trainY.shape))
```

#Visualizing Data

```
for i in range(6):
```

```
    plt.subplot(230+1+i)
```


```
    plt.imshow(trainX[i])
```

```
plt.show()
```

```
plt.imshow(trainX[88])
```


```
print("Label:", trainY[88])
```

→ Output :



Load and visualize MNIST

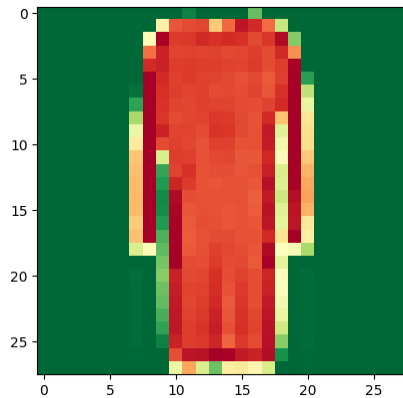
```
from keras.datasets import mnist
data = mnist.load_data()
(X_train, Y_train), (X_test, Y_test) = mnist.load_data()
plt.imshow(X_train[6], cmap='Greens')
plt.show()
print(Y_train[6])
```



Load and visualize Fashion MNIST

```
from keras.datasets import fashion_mnist
(x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()
```

```
plt.imshow(x_train[25], cmap='RdYlGn_r')
plt.show()
print("Label:", y_train[25])
```



Text Vectorization

```
Text_data = np.array([[ "I am Vikki"], [ "Visca Barca"], [ "I am a fan of Sports"]])
Text_data
```

→ **Output :**

```
array([[ 'Iam Vikki'],
       [ 'Visca Barca'],
       [ 'I am a fan of Sports']], dtype='<U20')

```

#Applying Text Vectorization

```
vectorizer = TextVectorization(output_mode='int')
vectorizer.adapt(Text_data)
integer_data = vectorizer(Text_data)
print("Vectorized Data:", integer_data.numpy())
print("Vocabulary:", vectorizer.get_vocabulary())
```

→ **Output :**

```
[', '[UNK]',
 np.str_('vikki'),
 np.str_('sports'),
 np.str_('of'),
 np.str_('iam'),
 np.str_('i'),
 np.str_('fan'),
 np.str_('Barca'),
 np.str_('am'),
 np.str_('Visca'),
 np.str_('a')]

```

Normalization

```
data = np.random.randint(0, 256, size=(64, 200, 200, 3)).astype('float32')
```

```
print("Original Mean:", np.mean(data))
```

→ **Output :** Mean: 127.53194

```
print("Original Variance:", np.var(data))
```

→ **Output :** Variance: 5459.2676

```
Normalizer = Normalization(axis=None)
```

```
Normalizer.adapt(data)
```

```
Normalized = Normalizer(data)
```

```
print("Normalized Mean:", np.mean(Normalized))
```

→ **Output :** Mean: 2.853652e-07

```
print("Normalized Variance:", np.var(Normalized))
```

→ **Output :** Variance: 0.9999998

Result:

The experiment successfully demonstrated key features of Keras such as dataset loading, image visualization, text vectorization, and data normalization using TensorFlow layers.

Experiment: 7 Logistic Regression

Aim:

To implement Logistic Regression using TensorFlow for multi-class classification on the Iris dataset.

Software Requirements:

- Google Colab

Program:

Importing libraries

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import tensorflow as tf
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
```

Load dataset

```
data = load_iris()
x = data["data"]
y = data["target"]
```

Create a DataFrame for reference

```
dataset = pd.DataFrame(data=np.concatenate((x, y.reshape(-1, 1)), axis=1),
                       columns=["Sepal length", "Sepal width", "Petal length", "Petal width", "target"])
print(dataset.head())
```

Split dataset

```
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.1, shuffle=True)
```


Build logistic regression model using Keras API

```
input = tf.keras.Input(shape=(4,))  
X = tf.keras.layers.Dense(3, activation='sigmoid')(input)  
model = tf.keras.models.Model(input, X)
```

Compile the model

```
model.compile(optimizer=tf.keras.optimizers.SGD(learning_rate=0.01),  
              loss=tf.keras.losses.sparse_categorical_crossentropy,  
              metrics=["accuracy"])
```

Train the model

```
train = model.fit(x_train, y_train, validation_data=(x_test, y_test), epochs=200)
```

Prediction for a new data point

```
New_data = [5.7, 2.8, 4.1, 1.3]  
y_pred = model.predict(np.array(New_data).reshape(1, -1))  
print("Prediction (Raw):", y_pred)  
print("Predicted Class:", np.argmax(y_pred))
```

Output:

```
predicted: [[0.0595241  0.31964937 0.28121248]]  
Predicted Class: 1
```

Result:

The logistic regression model was successfully implemented using TensorFlow. It accurately classified the Iris dataset into one of the three species.

Experiment: 8 Multi-Layer Perceptron

Aim:

To build and evaluate a Multi-Layer Perceptron (MLP) model using TensorFlow/Keras for classification of the Iris dataset.

Software Requirements:

- Google Colab

Program:

Importing libraries

```
from numpy import argmax
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
import numpy as np
import pandas as pd
```

Keras Model

```
from tensorflow.keras import Sequential
from tensorflow.keras.layers import Dense
```

Load Iris dataset

```
iris = load_iris()
x = iris.data.astype('float32')
y = iris.target
```

Create DataFrame

```
iris_df = pd.DataFrame(x, columns=iris.feature_names)
print("Iris Dataset Preview:")
print(iris_df.head())
```

Train-Test Split

```
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.25)
```

Model input feature size

```
n_features = x_train.shape[1]
print("Number of Features:", n_features)
```

Build MLP model

```
model = Sequential()
model.add(Dense(10, activation='relu', kernel_initializer='he_normal',
input_shape=(n_features,)))
model.add(Dense(4, activation='relu', kernel_initializer='he_normal'))
model.add(Dense(3, activation='softmax')) # 3 output classes
```

Compile the model

```
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
```

Train the model

```
model.fit(x_train, y_train, epochs=150, batch_size=32)
```

Evaluate the model

```
loss, accuracy = model.evaluate(x_test, y_test)
print(f"Test Accuracy: {accuracy * 100:.2f}%")
print(f"Test Loss: {loss:.4f}")
```

Prediction on new data

```
new_data = np.array([ [5.1, 3.5, 1.4, 0.2], [6.4, 3.1, 5.5, 1.8], [6.7, 3.3, 5.7, 2.5]
                      [4.6, 3.4, 1.4, 0.3],[4.8, 3.4, 1.9, 0.2],[5.6, 2.5, 3.9, 1.1]])
```

```
y_pred = model.predict(new_data)
print("Predictions (probabilities):")
print(y_pred)
print("Predicted Classes:")
```

```
print(np.argmax(y_pred, axis=1))
```

Output :

```
Test Accuracy: 100.00%
```

```
Test Loss: 0.0219
```

Output :

```
Model Evaluation Score: 1.0
```

Output :

```
Prediction: [[9.9630344e-01 3.6965355e-03 3.2302991e-13]
 [2.3607190e-09 3.3192713e-02 9.6680731e-01]
 [1.3617095e-10 7.7143717e-03 9.9228561e-01]
 [9.9582160e-01 4.1784509e-03 4.3446578e-13]
 [9.9535805e-01 4.6419362e-03 5.6031872e-13]
 [6.1574858e-04 9.9901319e-01 3.7104887e-04]]
Prediction Class: 16
```

Result:

The Multi-Layer Perceptron (MLP) model trained using the Iris dataset and successfully Verified.

Experiment: 9 Image Recognition CNN

Aim:

To develop an image classification model using Convolutional Neural Networks (CNN) on the Fashion MNIST dataset.

Software Requirements:

- Google Colab

Program:

Importing libraries

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow import keras
from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
```

Load the Fashion MNIST dataset

```
fashion = keras.datasets.fashion_mnist
(x_train_sp, y_train_sp), (x_test_sp, y_test_sp) = fashion.load_data()
```

Fashion item labels

```
item_names = ["T-Shirt/Top", "Trouser", "Pullover", "Dress", "Coat",
              "Sandal", "Shirt", "Sneaker", "Bag", "Ankle Boot"]
```

Reshape the dataset to add the channel dimension (1 for grayscale)

```
x_train_sp = x_train_sp.reshape((60000, 28, 28, 1))
x_test_sp = x_test_sp.reshape((10000, 28, 28, 1))
```

Normalize the pixel values to range 0-1

```
x_train_norm = x_train_sp / 255.0
x_test_norm = x_test_sp / 255.0
```

Split validation data from training data

```
x_validate, x_train = x_train_norm[:5000], x_train_norm[5000:]  
y_validate, y_train = y_train_sp[:5000], y_train_sp[5000:]  
x_test = x_test_norm
```

Set random seed for reproducibility

```
tf.random.set_seed(42)
```

Build CNN model

```
model = keras.models.Sequential([  
    keras.layers.Conv2D(filters=32, kernel_size=(3,3), activation="relu",  
input_shape=[28,28,1]),  
    keras.layers.MaxPooling2D(pool_size=(2,2)),  
    keras.layers.Flatten(),  
    keras.layers.Dense(300, activation="relu"),  
    keras.layers.Dense(100, activation="relu"),  
    keras.layers.Dense(10, activation="softmax")  
])
```

Display model summary

```
model.summary()
```

Compile the model

```
model.compile(optimizer='adam',  
    loss='sparse_categorical_crossentropy',  
    metrics=['accuracy'])
```

Train the model

```
history = model.fit(x_train, y_train, epochs=10, batch_size=32,  
    validation_data=(x_validate, y_validate))
```

Evaluate model on test data

```
test_loss, test_accuracy = model.evaluate(x_test, y_test_sp)
print(f"Test Accuracy: {test_accuracy * 100:.2f}%")
```

Predicting on test data

```
y_pred = model.predict(x_test)
y_pred_classes = np.argmax(y_pred, axis=1)
```

Confusion matrix and classification report

```
print("\nClassification Report:\n")
print(classification_report(y_test_sp, y_pred_classes, target_names=item_names))
```

→ Output :

```
accuracy: 0.91527 - loss: 0.2513
Test Accuracy: 89.56%
```

```
Classification Report:
              precision    recall  f1-score   support

T-Shirt/Top      0.40      0.87      0.55      1000
Trouser          0.00      0.00      0.00      1000
Pullover         0.31      0.75      0.44      1000
Dress            0.16      0.81      0.27      1000
Coat             0.00      0.00      0.00      1000
Sandal           0.00      0.00      0.00      1000
Shirt            0.00      0.00      0.00      1000
Sneakers         0.70      0.24      0.36      1000
Bag              0.00      0.00      0.00      1000
Ankle Boot       1.00      0.00      0.00      1000
accuracy         0.92      0.91      0.91     10000
macro avg        0.26      0.27      0.16     10000
weighted avg      0.26      0.27      0.16     10000
```

Result:

The CNN model was successfully built and trained using the Fashion MNIST dataset. It learned to recognize different clothing items and achieved good accuracy on the test data.

Experiment: 10 Transfer Learning for Audio Classification

Aim:

To perform audio classification using Transfer Learning with the YAMNet model, by extracting audio embeddings and training a custom classifier to classify animal sounds (dog and cat) from the ESC-50 dataset.

Software Requirements:

➤ Google Colab

Program:

Importing libraries

```
import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import tensorflow as tf
import tensorflow_hub as hub
import tensorflow_io as tfio
from IPython import display
```

Load YAMNet model from TensorFlow Hub

```
yamnet_model_handle = "https://tfhub.dev/google/yamnet/1"
yamnet_model = hub.load(yamnet_model_handle)
```

Load test audio

```
testing_wav_file_name = tf.keras.utils.get_file('miaow_16k.wav',
    'https://storage.googleapis.com/audioset/miaow_16k.wav',
    cache_dir='.', cache_subdir='test_data')
```

Function to load and resample audio

```
def load_wav_16k_mono(filename):
    file_contents = tf.io.read_file(filename)
```



```
wav, sample_rate = tf.audio.decode_wav(file_contents, desired_channels=1)
wav = tf.squeeze(wav, axis=-1)
sample_rate = tf.cast(sample_rate, dtype=tf.int64)
wav = tfio.audio.resample(wav, rate_in=sample_rate, rate_out=16000)
return wav

testing_wav_data = load_wav_16k_mono(testing_wav_file_name)
```

Play audio

```
display.Audio(testing_wav_data, rate=16000)
```

Load class names

```
class_map_path = yamnet_model.class_map_path().numpy().decode('utf-8')
class_names = list(pd.read_csv(class_map_path)['display_name'])
```

Run YAMNet prediction

```
scores, embeddings, spectrogram = yamnet_model(testing_wav_data)
class_scores = tf.reduce_mean(scores, axis=0)
top_class = tf.math.argmax(class_scores)
print(f"The main sound is: {class_names[top_class]}")
print(f"The embeddings shape: {embeddings.shape}")
```

Load ESC-50 dataset

```
_ = tf.keras.utils.get_file('ESC-50.zip',
    'https://github.com/karoldvl/ESC-50/archive/master.zip',
    cache_dir='.', cache_subdir='datasets', extract=True)
```

Read metadata and filter dog and cat

```
esc50_csv = './datasets/ESC-50-master/meta/esc50.csv'
base_data_path = './datasets/ESC-50-master/audio/'
pd_data = pd.read_csv(esc50_csv)
```

```
my_classes = ['dog', 'cat']
map_class_to_id = {'dog': 0, 'cat': 1}
filtered_pd = pd_data[pd_data.category.isin(my_classes)]
filtered_pd = filtered_pd.assign(target=filtered_pd['category'].map(map_class_to_id))
filtered_pd['filename'] = filtered_pd['filename'].apply(lambda name:
os.path.join(base_data_path, name))
```

Prepare dataset

```
filenames = filtered_pd['filename']
targets = filtered_pd['target']
folds = filtered_pd['fold']
main_ds = tf.data.Dataset.from_tensor_slices((filenames, targets, folds))

def load_wav_for_map(filename, label, fold):
    return load_wav_16k_mono(filename), label, fold

main_ds = main_ds.map(load_wav_for_map)
```

Extract embeddings from YAMNet

```
def extract_embedding(wav_data, label, fold):
    scores, embeddings, spectrogram = yamnet_model(wav_data)
    num_embeddings = tf.shape(embeddings)[0]
    return embeddings, tf.repeat(label, num_embeddings), tf.repeat(fold, num_embeddings)

main_ds = main_ds.map(extract_embedding).unbatch()
cached_ds = main_ds.cache()
```

Split dataset

```
train_ds = cached_ds.filter(lambda emb, label, fold: fold < 4).map(lambda emb, label, fold:
(emb, label))
val_ds = cached_ds.filter(lambda emb, label, fold: fold == 4).map(lambda emb, label, fold:
(emb, label))
test_ds = cached_ds.filter(lambda emb, label, fold: fold == 5).map(lambda emb, label, fold:
```

```
(emb, label))
```

Batch and prefetch

```
train_ds = train_ds.shuffle(1000).batch(32).prefetch(tf.data.AUTOTUNE)
```

```
val_ds = val_ds.batch(32).prefetch(tf.data.AUTOTUNE)
```

```
test_ds = test_ds.batch(32).prefetch(tf.data.AUTOTUNE)
```

Build and train model

```
my_model = tf.keras.Sequential([
```

```
    tf.keras.layers.Input(shape=(1024,), name='input_embedding'),
```

```
    tf.keras.layers.Dense(512, activation='relu'),
```

```
    tf.keras.layers.Dense(len(my_classes))
```

```
], name='my_model')
```

```
my_model.compile(loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
```

```
                  optimizer='adam',
```

```
                  metrics=['accuracy'])
```

```
callback = tf.keras.callbacks.EarlyStopping(monitor='loss', patience=5,  
restore_best_weights=True)
```

```
history = my_model.fit(train_ds, epochs=20, validation_data=val_ds, callbacks=[callback])
```

Evaluate on test set

```
loss, accuracy = my_model.evaluate(test_ds)
```

```
print("Loss:", loss)
```

```
print("Accuracy:", accuracy)
```

Predict new sound class



```
scores, embeddings, spectrogram = yamnet_model(testing_wav_data)
```

```
result = my_model(embeddings).numpy() ==
```

```
inferred_class = my_classes[result.mean(axis=0).argmax()]
```

```
print(f"The main sound is: {inferred_class}")
```

→ Output :



Loss: 0.3009394705295563

Accuracy: 0.893750011920929

The main sound is: cat

Result:

The audio classification model was successfully built using transfer learning with YAMNet.

Experiment: 11 Stock Prediction Using LSTM

Aim:

To develop a deep learning model using LSTM (Long Short-Term Memory) for predicting Tesla stock prices based on historical data.

Software Requirements:

- Google Colab

Program:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

from sklearn.preprocessing import MinMaxScaler
from keras.models import Sequential
from keras.layers import LSTM, Dropout, Dense

# Load the dataset
tesla = pd.read_csv("/content/drive/MyDrive/Colab Notebooks/Deep
Learning/LSTM(RNN)/Tesla DataSet 5 years.csv")
print(tesla.head())

# Data splitting
training = tesla.iloc[:800, 1:2].values # "Open" prices
testing = tesla.iloc[800:, 1:2].values # Remaining for testing

# Data normalization
nm_scale = MinMaxScaler(feature_range=(0,1))
training_scaled = nm_scale.fit_transform(training)

# Create data structure with 60 time steps
```

```
x_train = []
y_train = []
for i in range(60, 800):
    x_train.append(training_scaled[i-60:i, 0])
    y_train.append(training_scaled[i, 0])

x_train, y_train = np.array(x_train), np.array(y_train)
x_train = np.reshape(x_train, (x_train.shape[0], x_train.shape[1], 1)) # 3D input for LSTM
```

Model architecture

```
model = Sequential()
model.add(LSTM(units=50, return_sequences=True, input_shape=(x_train.shape[1], 1)))
model.add(Dropout(0.2))
model.add(LSTM(units=50, return_sequences=True))
model.add(Dropout(0.2))
model.add(LSTM(units=50, return_sequences=True))
model.add(Dropout(0.2))
model.add(LSTM(units=50))
model.add(Dropout(0.2))
model.add(Dense(units=1))
```

Compile and train the model

```
model.compile(optimizer='adam', loss='mean_squared_error')
model.fit(x_train, y_train, epochs=100, batch_size=32)
```

Preparing test inputs

```
total_data = pd.concat((tesla.iloc[:800, 1:2], tesla.iloc[800:, 1:2]), axis=0)
inputs = total_data[len(total_data) - len(testing) - 60:].values
inputs = inputs.reshape(-1, 1)
```

```
inputs = nm_scale.transform(inputs)
```

```
x_test = []
```

```
for i in range(60, 60 + len(testing)):
```

```
    x_test.append(inputs[i-60:i, 0])
```

```
x_test = np.array(x_test)
```

```
x_test = np.reshape(x_test, (x_test.shape[0], x_test.shape[1], 1))
```

Predicting the stock prices

```
predicted_stock_price = model.predict(x_test)
```

```
predicted_stock_price = nm_scale.inverse_transform(predicted_stock_price)
```

Actual stock prices

```
actual_stock_price = testing
```

Plotting results

```
plt.figure(figsize=(12,6))
```

```
plt.plot(tesla.loc[800:, "Date"], actual_stock_price, color="red", label="Actual Tesla Stock Price")
```

```
plt.plot(tesla.loc[800:, "Date"], predicted_stock_price, color="green", label="Predicted Tesla Stock Price")
```

```
plt.title("Tesla Stock Price Prediction")
```

```
plt.xlabel("Time")
```

```
plt.ylabel("Tesla Stock Price")
```


```
plt.xticks(rotation=45)
```

```
plt.legend()
```

```
plt.tight_layout()
```

```
plt.show()
```

→ Output :



Result:

The LSTM-based Recurrent Neural Network was trained using Tesla stock data and successfully predicted future stock prices with reasonable accuracy.

Experiment: 12 Image Denoising using Autoencoder

Aim:

To develop and train an Autoencoder model for denoising images using the Fashion MNIST dataset.

Software Requirements:

- Google Colab

Program:

```
import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.datasets import fashion_mnist
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input, Dense, Conv2D, MaxPooling2D, UpSampling2D
from tensorflow.keras.optimizers import Adam
```

Load Fashion MNIST dataset

```
(x_train, _), (x_test, _) = fashion_mnist.load_data()
```

Normalize & Reshape images

```
x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0
x_train = np.expand_dims(x_train, -1)
x_test = np.expand_dims(x_test, -1)
```

Add Noise

```
noise_factor = 0.5
x_train_noisy = x_train + noise_factor * np.random.normal(loc=0.0, scale=1.0,
size=x_train.shape)
x_test_noisy = x_test + noise_factor * np.random.normal(loc=0.0, scale=1.0,
size=x_test.shape)
x_train_noisy = np.clip(x_train_noisy, 0., 1.)
x_test_noisy = np.clip(x_test_noisy, 0., 1.)
```

Build Autoencoder

```
input_img = Input(shape=(28, 28, 1))
x = Conv2D(32, (3, 3), activation='relu', padding='same')(input_img)
x = MaxPooling2D((2, 2), padding='same')(x)
x = Conv2D(32, (3, 3), activation='relu', padding='same')(x)
encoded = MaxPooling2D((2, 2), padding='same')(x)
x = Conv2D(32, (3, 3), activation='relu', padding='same')(encoded)
x = UpSampling2D((2, 2))(x)
```

```
x = Conv2D(32, (3, 3), activation='relu', padding='same')(x)
x = UpSampling2D((2, 2))(x)
decoded = Conv2D(1, (3, 3), activation='sigmoid', padding='same')(x)
```

Compile and train

```
autoencoder = Model(input_img, decoded)
autoencoder.compile(optimizer=Adam(), loss='binary_crossentropy')
autoencoder.fit(x_train_noisy, x_train, epochs=10, batch_size=128, shuffle=True,
validation_data=(x_test_noisy, x_test))
```

Evaluate

```
decoded_imgs = autoencoder.predict(x_test_noisy)
```

Visualization

```
n = 10
plt.figure(figsize=(20, 4))
for i in range(n):
    ax = plt.subplot(2, n, i + 1)
    plt.imshow(x_test_noisy[i].reshape(28, 28), cmap='gray')
    plt.title("Noisy")
    plt.axis('off')

    ax = plt.subplot(2, n, i + 1 + n)
    plt.imshow(decoded_imgs[i].reshape(28, 28), cmap='gray')
    plt.title("Denoised")
    plt.axis('off')
plt.show()
```

Accuracy estimation


```
loss = autoencoder.evaluate(x_test_noisy, x_test, verbose=0)
print(f"Test Loss: {loss}")

decoded_imgs = autoencoder.predict(x_test_noisy)
diff = np.abs(decoded_imgs - x_test)
threshold = 0.1
correct_predictions = np.sum(diff < threshold)
accuracy = correct_predictions / (x_test.shape[0] * x_test.shape[1] * x_test.shape[2])
print(f"Accuracy based on threshold: {accuracy}")
```

→ Output :

```
Test Loss: 0.29475435614585876
```

```
Accuracy based on threshold: 0.7584261479591837
```



Result:

The Autoencoder model was successfully trained on the Fashion MNIST dataset and effectively denoised the noisy images.

Experiment: 13 Neural Style Transfer GAN

Aim:

To apply neural style transfer using a pre-trained GAN model from TensorFlow Hub to blend the style of one image with the content of another.

Software Requirements:

- Google Colab

Program:

```
from matplotlib import gridspec
import matplotlib.pyplot as plt
import numpy as np
import tensorflow as tf
import tensorflow_hub as tensorflow_hub
import PIL
from google.colab import files

# Image Processing Function
def load_image(image, image_size=(512,512)):
    img = tf.io.read_file(image)
    img = tf.image.decode_image(img, channels=3)
    img = tf.image.convert_image_dtype(img, tf.float32)
    img = tf.image.resize(img, image_size, preserve_aspect_ratio=True)
    img = img[tf.newaxis, :]
    return img

# Upload content image
upload = files.upload()
original_image = load_image("Lion.jpg")
```

Upload style image

```
upload = files.upload()  
style_image = load_image("Tiger.jpeg")
```

Display function

```
def imshow(image, title=None):  
    if len(image.shape) > 3:  
        image = tf.squeeze(image, axis=0)  
    plt.imshow(image)  
    if title:  
        plt.title(title)
```

Display content and style images

```
imshow(original_image, "Original Image")
```



```
imshow(style_image, "Style Image")
```



```
# Load model from TensorFlow Hub
```

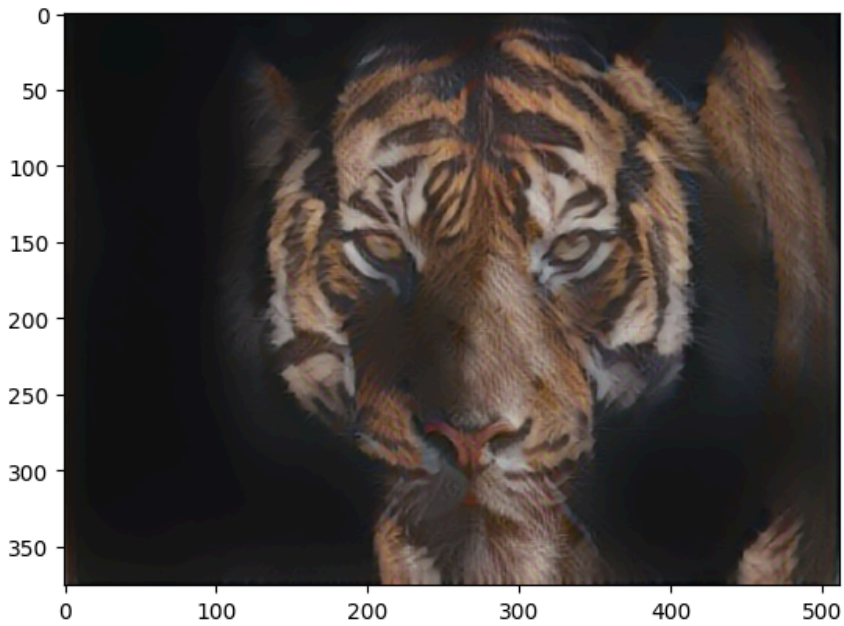
```
hub_handle = "https://tfhub.dev/google/magenta/arbitrary-image-stylization-v1-256/2"
```

```
hub_module = tensorflow_hub.load(hub_handle)
```

```
# Apply style transfer
```


```
outputs = hub_module(original_image, style_image)
```

```
imshow(outputs[0], "Stylized Image")
```



Optionally reverse style and content

```
outputs = hub_module(style_image, original_image)  
imshow(outputs[0], "Reversed Style Image")
```



Result:

The neural style transfer model was successfully implemented. The original image was stylized using the features of the style image, and the resulting images were visualized effectively.